

mnist_unet

November 1, 2018

```
In [ ]: import tensorflow as tf
import numpy as np

import innvestigate
import innvestigate.utils as iutils
from innvestigate.utils.visualizations import heatmap

from keras.models import Model
from keras.layers import concatenate, Input, Conv2D, MaxPooling2D, Conv2DTranspose, Activation
from keras.layers.normalization import BatchNormalization

from tqdm import tqdm_notebook as tqdm

import matplotlib.pyplot as plot
import numpy as np
import os
%matplotlib inline

In [ ]: from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets('MNIST_DATA', one_hot=True)

In [41]: def ConvBNRelu(input_tensor, filters):
    out = Conv2D(filters, (3, 3), activation=None, padding='valid', use_bias=False)(input_tensor)
    out = BatchNormalization(axis=-1)(out)
    return Activation('relu')(out)

def get_model(weights=None, flatten_output=False):
    filters = 4
    inputs = Input((28, 28, 1), name="input")

    conv1 = ConvBNRelu(inputs, filters)
    conv2 = ConvBNRelu(conv1, filters)

    pool1 = MaxPooling2D(pool_size=(2, 2))(conv2)

    conv3 = ConvBNRelu(pool1, 2*filters)
    conv4 = ConvBNRelu(conv3, 2*filters)
```

```

up1 = Conv2DTranspose(filters, (2, 2), strides=(2, 2), activation='relu', padding='
up1 = concatenate([up1, Cropping2D(cropping=(4, 4))(conv2)])

conv5 = ConvBNRelu(up1, filters)
conv6 = ConvBNRelu(conv5, filters)

logits = Conv2D(1, (1, 1), activation="sigmoid")(conv6)

if flatten_output:
    logits = Flatten()(logits)

model = Model(inputs=[inputs], outputs=[logits])
model.compile(loss="binary_crossentropy", optimizer="adam")
if weights is not None:
    model.load_weights(weights)
return model

```

In [43]: model = get_model()

```

batch_size = 32

x = mnist.train.images
count = len(x)
x = np.reshape(x, (count, 28, 28, 1))
y = np.where(x > 0.5, 1, 0)
y = y[:, 8:-8, 8:-8, :]
model.fit(x, y, batch_size=batch_size, epochs=2)
model.save_weights("params")

```

Epoch 1/2

55000/55000 [=====] - 66s 1ms/step - loss: 0.1222

Epoch 2/2

55000/55000 [=====] - 58s 1ms/step - loss: 0.0526

In [53]: model = get_model("params")

```



sample = x[0:1]
plot.imshow(sample.squeeze(), cmap="gray")
plot.show()


plot.imshow(sample[:, 8:-8, 8:-8].squeeze(), cmap="gray")
plot.show()

p = model.predict(sample)
p = p.squeeze()
#p = np.pad(p, (8, 8), mode="constant", constant_values=0)

```


```
i, j = np.unravel_index(np.argmax(p), p.shape)
plot.imshow(p.squeeze(), cmap="gray")
plot.show()
```






```
In [57]: pixel_x = 11
pixel_y = 10
model = get_model("params")
p = model.predict(sample)
print('Neuron output at (%d, %d) = ' % (pixel_y, pixel_x), p[0, pixel_y, pixel_x, 0])
index = np.ravel_multi_index((0, pixel_y, pixel_x, 0), p.shape)
flatten_model = get_model("params", flatten_output=True)
flat_p = flatten_model.predict(sample)
print('Flattened model neuron output at (%d, %d) = ' % (pixel_y, pixel_x), flat_p[0, index])
for _ in range(3):
    analyzer = innvestigate.analyzer.LRPEpsilon(flatten_model, input_layer_rule=(0, 1),
    analysis = analyzer.analyze(sample, index)
    print("relevance at index = ", analysis[0, pixel_y+8, pixel_x+8, 0])
    print("analysis.sum() = ", analysis.sum())
    #analysis = heatmap(analysis)
    plot.imshow(analysis.squeeze(), cmap='gray', interpolation='nearest')
    plot.show()
```


```
Neuron output at (10, 11) = 0.99998176
Flattened model neuron output at (10, 11) = 0.99998176
relevance at index = -8.912141
analysis.sum() = 3.5663242
```



```
relevance at index = -0.8532734  
analysis.sum() = 3.5663233
```



```
relevance at index = 0.7234432  
analysis.sum() = 3.5663233
```



```
In [ ]:
```